


Project Report: Restaurant Recommendation Chatbot Powered by AI

1. Objective

The primary objective of this project is to design and implement an **AI-powered chatbot** that recommends restaurants based on user preferences. The chatbot leverages natural language queries (e.g., “*Find me vegetarian restaurants under ₹500 in Indiranagar*”) and provides tailored recommendations using structured restaurant datasets combined with OpenAI’s language model capabilities.

2. Design



System Architecture

The chatbot system is structured as follows:

1. User Input Layer

- Users interact with the chatbot through natural language queries.

2. Preprocessing & Data Layer

- A dataset (`restaurant_dataset.csv`) containing restaurant information (cuisine, price, rating, location, etc.) is loaded and processed with **Pandas**.

3. AI Model Integration

- The **OpenAI Chat Completion API** is used to handle conversation, extract intent, and generate human-like responses.
- The system message defines the chatbot's role: "*You are a restaurant recommendation assistant.*"

4. Recommendation Engine

- Filters restaurants based on extracted constraints such as cuisine, budget, and location.
- Formats recommendations to present to the user.

5. Output Layer

- Chatbot returns clear, conversational responses containing a list of restaurants matching the request.
-

3. Implementation

3.1 Importing Required Libraries

- `pandas`: Used to load and manipulate the restaurant dataset.
- `openai`: To connect with the OpenAI Chat Completion API and generate chatbot responses.
- `json`: For structured data representation when passing responses.

This ensures the environment is set up for both **data handling** and **AI-powered conversation**.

3.2 Loading the Dataset

- A CSV file containing restaurant information was loaded using `pandas.read_csv()`.

- Columns included attributes such as:
 - **Name** (restaurant name)
 - **Cuisine** (e.g., Italian, Chinese, Indian)
 - **Location** (area/city)
 - **Price** (average cost for two people)
 - **Rating** (user ratings out of 5)

This provides a structured **knowledge base** for the chatbot.

3.3 Exploring and Cleaning Data

- Inspected the dataset (`df.head()`, `df.info()`, `df.describe()`) to understand structure.
- Checked for **null values or inconsistencies**.
- Converted certain columns to consistent formats (e.g., converting all cuisines to lowercase, ensuring numeric fields like price are integers).

This makes filtering reliable later on.

3.4 Filtering Logic (Recommendation Engine)

- Functions were written to filter restaurants based on user constraints such as:
 - **Cuisine** (e.g., “Chinese”)
 - **Location** (e.g., “Indiranagar”)
 - **Price range** (e.g., under ₹500)
 - **Ratings** (e.g., >4.0)

This acts as the **retrieval engine**.

3.5 Integrating with OpenAI API

- OpenAI's **Chat Completions API** (`gpt-3.5-turbo`) was integrated.
 - Chatbot roles were defined:
 - **System role:** Sets context → “You are a restaurant recommendation assistant.”
 - **User role:** User’s input query.
 - **Assistant role:** LLM generates conversational response.
- This ensures the **chatbot understands natural language**.

3.6 Combining Dataset Lookup + AI Responses

- When the user asks, the chatbot:
 1. Extracts intent (cuisine, budget, location).
 2. Queries the dataset using the filtering function.
 3. Passes the filtered results to OpenAI to generate a **well-phrased conversational response**.

4. Challenges

1. **Dataset Limitations**
 - The dataset may lack certain fields (e.g., ambience, parking, dietary preferences). This limits the richness of recommendations.
2. **Handling Ambiguity in User Input**
 - Natural language inputs like “*a cheap cozy place for date night*” are hard to map directly to structured dataset filters.
3. **API Dependency**
 - The chatbot depends heavily on OpenAI API availability and cost, which could be limiting in production-scale scenarios.

4. Scalability

- Current implementation works for small datasets but may not handle very large datasets efficiently without optimization (e.g., indexing, embeddings).
-

5. Lessons Learned

1. Combining AI with Structured Data

- Pure LLM-based chatbots are good at conversation, but grounding them in **real datasets** makes recommendations accurate and reliable.

2. Importance of Context Management

- Defining system prompts carefully (e.g., “You are a restaurant recommendation assistant”) helps keep the chatbot focused on its domain.

3. Need for Moderation & Intent Filtering

- Moderation ensures safe interaction, while intent filtering makes sure the chatbot stays relevant to restaurant recommendations.

4. Future Enhancements

- Adding personalization (user history).
- Using embeddings for semantic search (e.g., “romantic rooftop places”).
- Expanding dataset fields for richer filtering (dietary restrictions, ambience, services).